

# Distributed Systems

ECE428

Lecture 5

*Adopted from Spring 2021*

# Recap: Timestamps Summary

- Comparing timestamps across events is useful.
  - Reconciling updates made to an object in a distributed datastore.
  - Rollback recovery during failures:
    1. Checkpoint state of the system;
    2. Log events (with timestamps);
    3. Rollback to checkpoint and replay events in order if system crashes.
- How to compare timestamps across different processes?
  - Physical timestamp: requires clock synchronization.
    - Google's Spanner Distributed Database uses "TrueTime".
  - Lamport's timestamps: cannot fully differentiate between causal and concurrent ordering of events.
    - Oracle uses "System Change Numbers" based on Lamport's clock.
  - Vector timestamps: larger message sizes.
    - Amazon's DynamoDB uses vector clocks.

# Recap: Timestamps Summary

- Comparing timestamps across events is useful.
  - Reconciling updates made to an object in a distributed datastore.
  - Rollback recovery during failures:
    - 1. Checkpoint state of the system; 2. Log events (with timestamps); 3. Rollback to checkpoint and replay events in order if system crashes.*
- How to compare timestamps across different processes?
  - Physical timestamp: requires clock synchronization.
    - Google's Spanner Distributed Database uses "TrueTime".
  - Lamport's timestamps: cannot fully differentiate between causal and concurrent ordering of events.
    - Oracle uses "System Change Numbers" based on Lamport's clock.
  - Vector timestamps: larger message sizes.
    - Amazon's DynamoDB uses vector clocks.

# Today's agenda

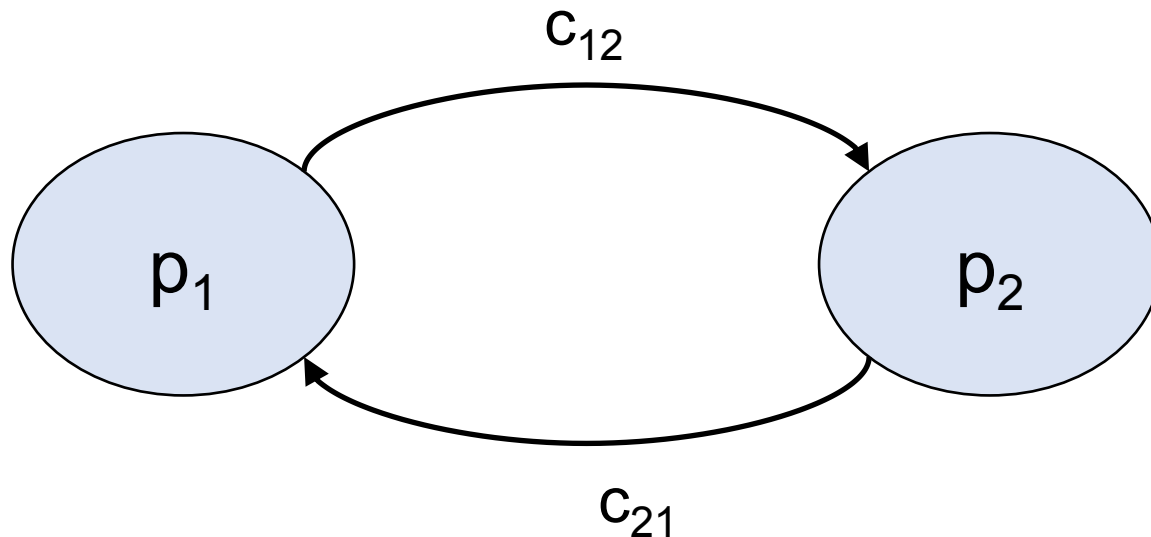
- Global State
  - Chapter 14.5
  - Goal: reason about how to capture the state across all processes of a distributed system without requiring time synchronization.

# Process, state, events

- Consider a system with  $n$  processes:  $\langle p_1, p_2, p_3, \dots, p_n \rangle$ .
- Each process  $p_i$  is associated with *state*  $s_i$ .
  - State includes values of all local variables, affected files, etc.
- Each channel can also be associated with a state.
  - Which messages are currently *pending* on the channel.
  - Can be computed from process' state:
    - Record when a process sends and receives messages.
    - if  $p_i$  sends a message that  $p_j$  has not yet received, it is pending on the channel.
- State of a process (or a channel) gets transformed when an *event occurs*. 3 types of events:
  - local computation, sending a message, receiving a message.

# Global State (or Global Snapshot)

- State of each process (and each channel) in the system at a given instant of time.
- Example:



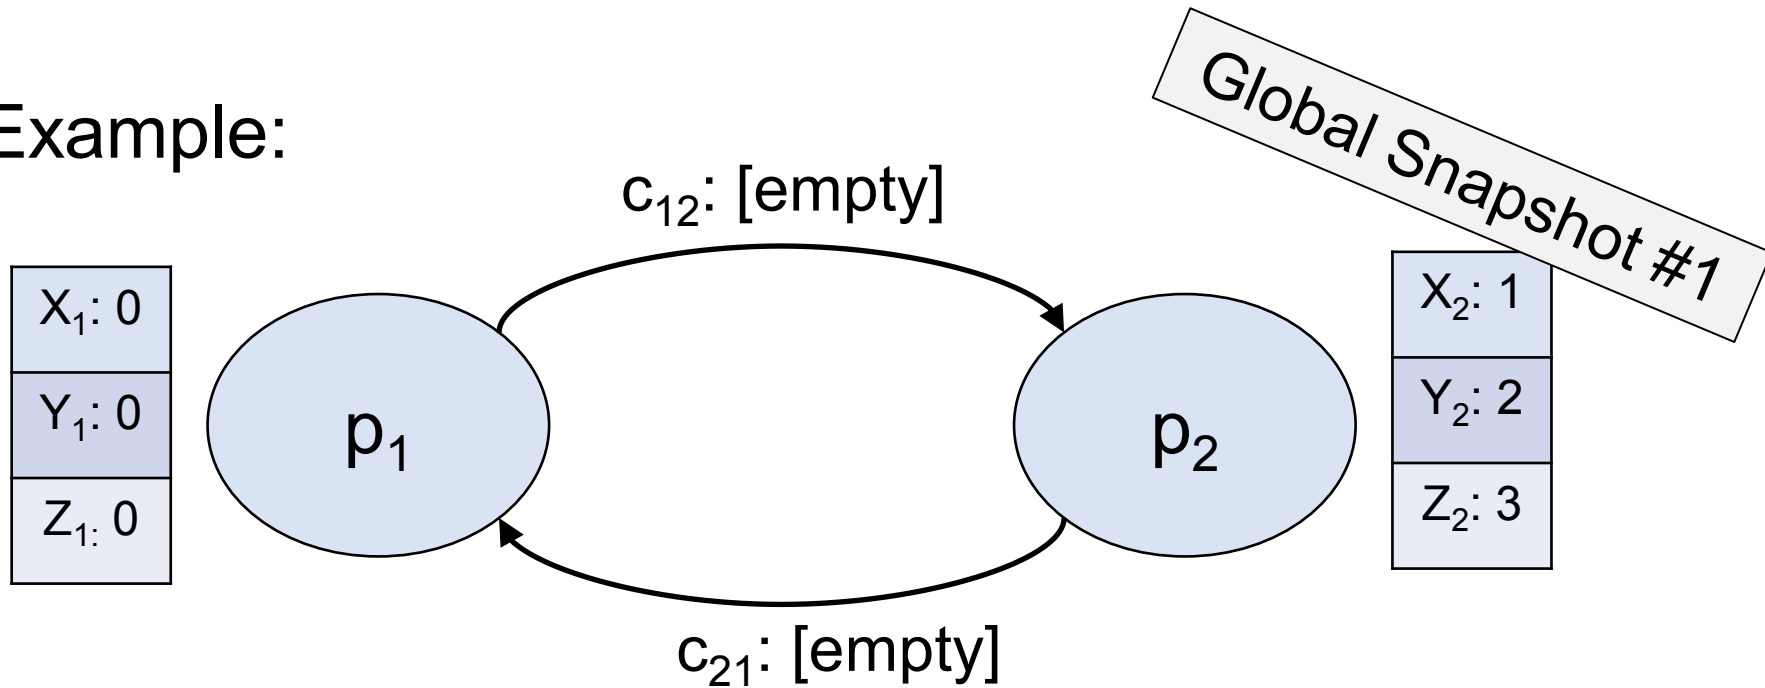
Two processes:  $p_1$  and  $p_2$ .

$c_{12}$ : channel from  $p_1$  to  $p_2$ .

$c_{21}$ : channel from  $p_2$  to  $p_1$ .

# Global State (or Global Snapshot)

- State of each process (and each channel) in the system at a given instant of time.
- Example:




Process state for  $p_1$  and  $p_2$ .  
No pending messages on the channels.

# Global State (or Global Snapshot)

- State of each process (and each channel) in the system at a given instant of time.

- Example:



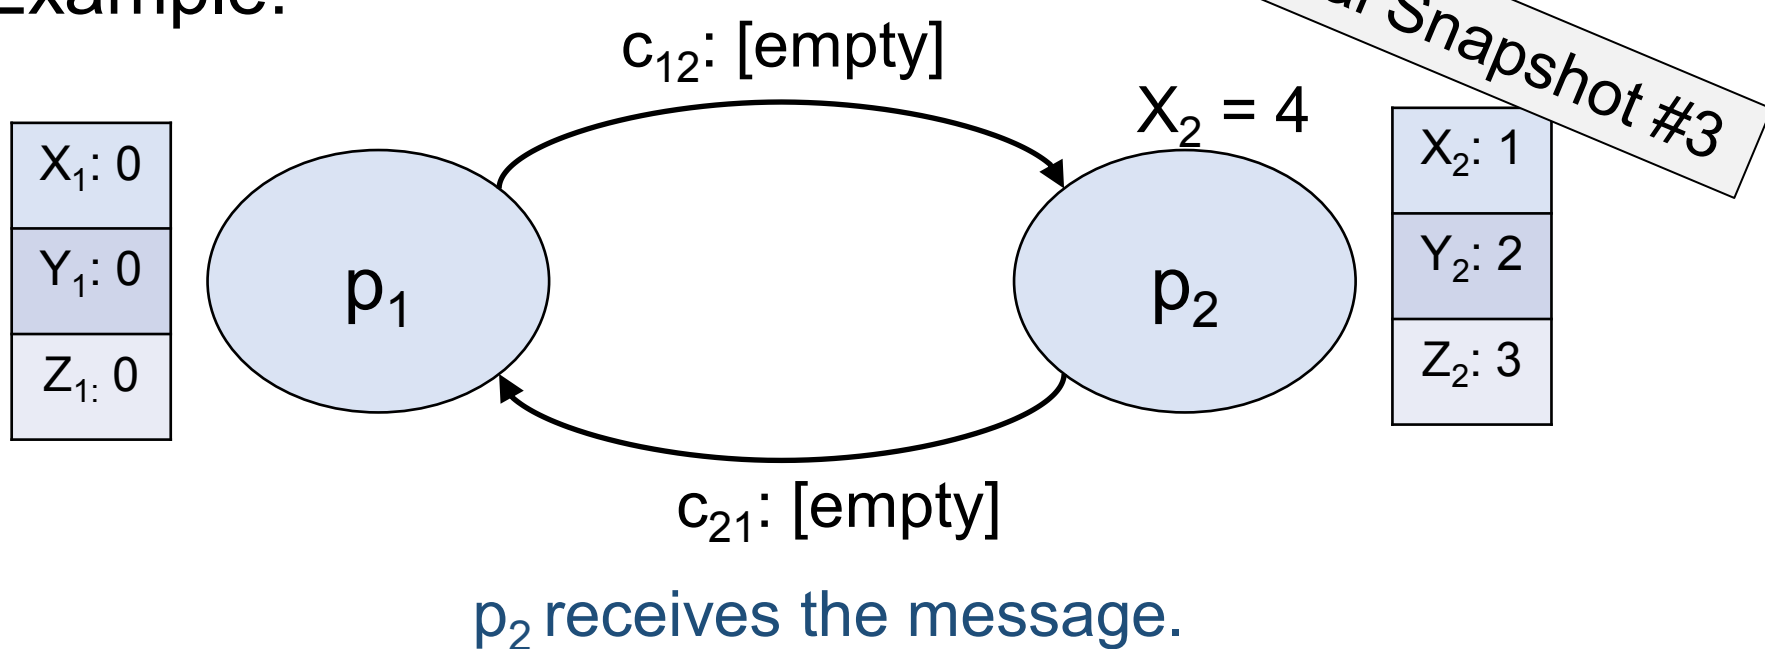
$p_1$  send a message to  $p_2$  asking it to set  $X_2 = 4$



# Global State (or Global Snapshot)

- State of each process (and each channel) in the system at a given instant of time.

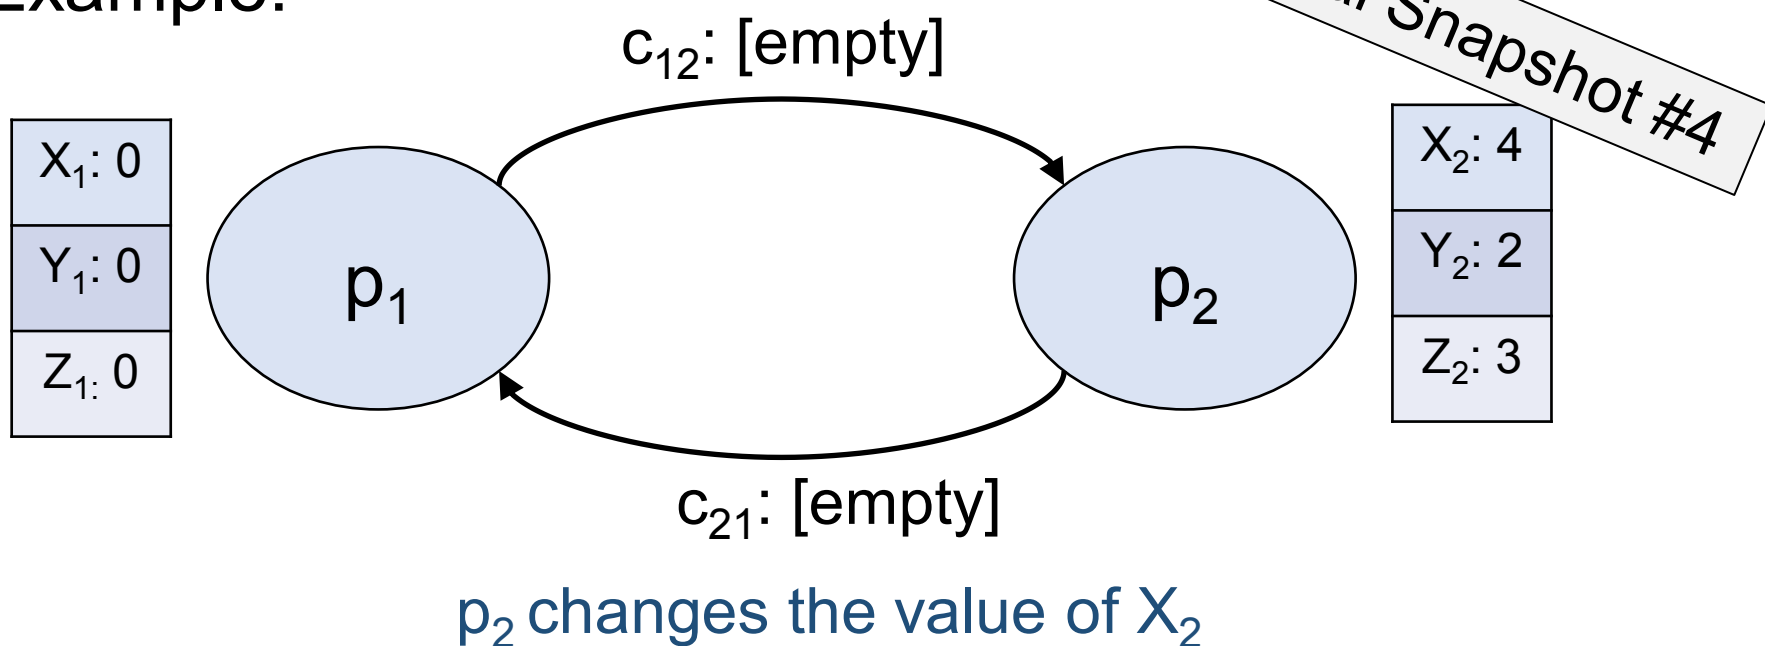
- Example:



# Global State (or Global Snapshot)

- State of each process (and each channel) in the system at a given instant of time.

- Example:



# Capturing a global snapshot

- Useful to capture a global snapshot of the system:
  - *Checkpointing* the system state.
  - Reasoning about unreferenced objects (for garbage collection).
  - Deadlock detection.
  - Distributed debugging.

# Capturing a global snapshot

- Difficult to capture a global snapshot of the system.
- Global state or global snapshot is state of each process (and each channel) in the system at a given *instant of time*.
- Strawman:
  - Each process records its state at 3:15pm.
  - We get the global state of the system at 3:15pm.
  - *But precise clock synchronization is difficult to achieve.*
- How do we capture global snapshots without precise time synchronization across processes?

# More notations and definitions

- State of a process (or a channel) gets transformed when an *event* occurs.
- 3 types of events:
  - local computation, sending a message, receiving a message.
- $e_i^n$  is the  $n^{\text{th}}$  event at  $p_i$ .

# More notations and definitions

- For a process  $p_i$ , where events  $e_i^0, e_i^1, \dots$  occur:

**history**( $p_i$ ) =  $h_i = \langle e_i^0, e_i^1, \dots \rangle$

**prefix history**( $p_i^k$ ) =  $h_i^k = \langle e_i^0, e_i^1, \dots, e_i^k \rangle$

**$s_i^k$**  :  $p_i$ 's state immediately after  $k^{\text{th}}$  event.

- For a set of processes  $\langle p_1, p_2, p_3, \dots, p_n \rangle$ :

**global history**:  $H = \cup_i (h_i)$

**global state**:  $S = \cup_i (s_i)$

# More notations and definitions

- For a process  $p_i$ , where events  $e_i^0, e_i^1, \dots$  occur:

**history**( $p_i$ ) =  $h_i = \langle e_i^0, e_i^1, \dots \rangle$

**prefix history**( $p_i^k$ ) =  $h_i^k = \langle e_i^0, e_i^1, \dots, e_i^k \rangle$

**$s_i^k$**  :  $p_i$ 's state immediately after  $k^{\text{th}}$  event.

- For a set of processes  $\langle p_1, p_2, p_3, \dots, p_n \rangle$ :

**global history**:  $H = \cup_i (h_i)$

**global state**:  $S = \cup_i (s_i)$

*But state at what time instant?*

# More notations and definitions

- For a process  $p_i$ , where events  $e_i^0, e_i^1, \dots$  occur:

**history**( $p_i$ ) =  $h_i = \langle e_i^0, e_i^1, \dots \rangle$

**prefix history**( $p_i^k$ ) =  $h_i^k = \langle e_i^0, e_i^1, \dots, e_i^k \rangle$

$s_i^k$ :  $p_i$ 's state immediately after  $k^{\text{th}}$  event.

- For a set of processes  $\langle p_1, p_2, p_3, \dots, p_n \rangle$ :

**global history**:  $H = \cup_i (h_i)$

**global state**:  $S = \cup_i (s_i^{k_i})$


a **cut**  $C \subseteq H = h_1^{c_1} \cup h_2^{c_2} \cup \dots \cup h_n^{c_n}$

the **frontier** of  $C = \{e_i^{c_i}, i = 1, 2, \dots, n\}$

**global state**  $S$  that corresponds to cut  $C = \cup_i (s_i^{c_i})$



# Example: Cut



$$C_A: \langle e_1^0, e_2^0 \rangle$$

Frontier of  $C_A$ :

$$C_B: \langle e_1^0, e_1^1, e_1^2, e_2^0, e_2^1, e_2^2 \rangle$$

Frontier of  $C_B$ :

# More notations and definitions

- For a process  $p_i$ , where events  $e_i^0, e_i^1, \dots$  occur:

**history**( $p_i$ ) =  $h_i = \langle e_i^0, e_i^1, \dots \rangle$

**prefix history**( $p_i^k$ ) =  $h_i^k = \langle e_i^0, e_i^1, \dots, e_i^k \rangle$

$s_i^k$ :  $p_i$ 's state immediately after  $k^{\text{th}}$  event.

- For a set of processes  $\langle p_1, p_2, p_3, \dots, p_n \rangle$ :

**global history**:  $H = \cup_i (h_i)$

**global state**:  $S = \cup_i (s_i^{k_i})$

a **cut**  $C \subseteq H = h_1^{c_1} \cup h_2^{c_2} \cup \dots \cup h_n^{c_n}$

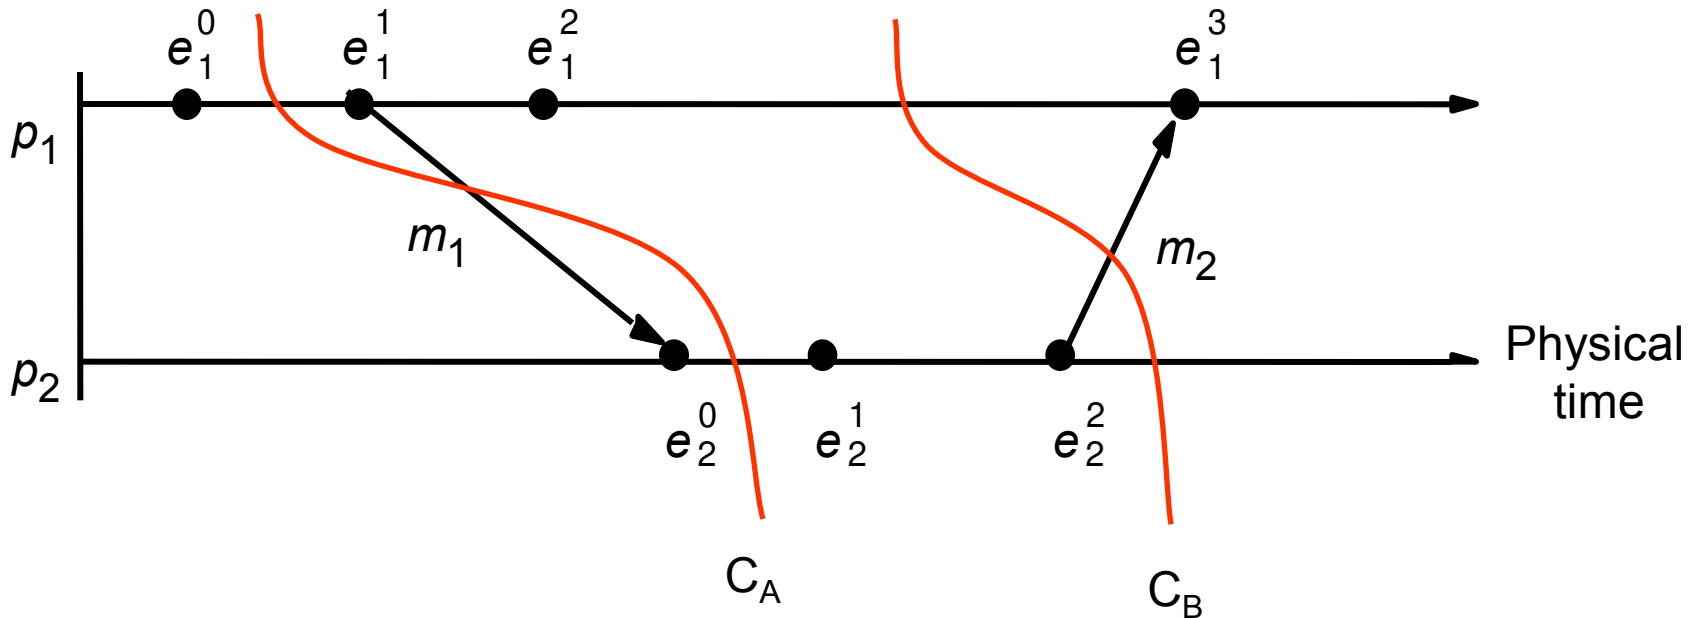
the **frontier** of  $C = \{e_i^{c_i}, i = 1, 2, \dots, n\}$

**global state**  $S$  that corresponds to cut  $C = \cup_i (s_i^{c_i})$

# Consistent cuts and snapshots

- A cut  $C$  is **consistent** if and only if
$$\forall e \in C \text{ (if } f \rightarrow e \text{ then } f \in C)$$

# Example: Cut



$$C_A: \langle e_1^0, e_2^0 \rangle$$

Frontier of  $C_A$ :  $\{e_1^0, e_2^0\}$

**Inconsistent cut.**

$$C_B: \langle e_1^0, e_1^1, e_1^2, e_2^0, e_2^1, e_2^2 \rangle$$

Frontier of  $C_B$ :  $\{e_1^2, e_2^2\}$

**Consistent cut.**

# Consistent cuts and snapshots

- A cut  $C$  is **consistent** if and only if
$$\forall e \in C \text{ (if } f \rightarrow e \text{ then } f \in C)$$
- A global state  $S$  is consistent if and only if it corresponds to a consistent cut.

# Consistent cuts and snapshots

- A cut  $C$  is **consistent** if and only if
$$\forall e \in C \text{ (if } f \rightarrow e \text{ then } f \in C)$$
- A global state  $S$  is consistent if and only if it corresponds to a consistent cut.
- *How do we find consistent global states?*

# Chandy-Lamport Algorithm

- Goal:
  - Record a global snapshot
    - Process state (and channel state) for a set of processes.
  - The recorded global state is consistent.
- Identifies a consistent cut.
- Records corresponding state locally at each process.

# Chandy-Lamport Algorithm

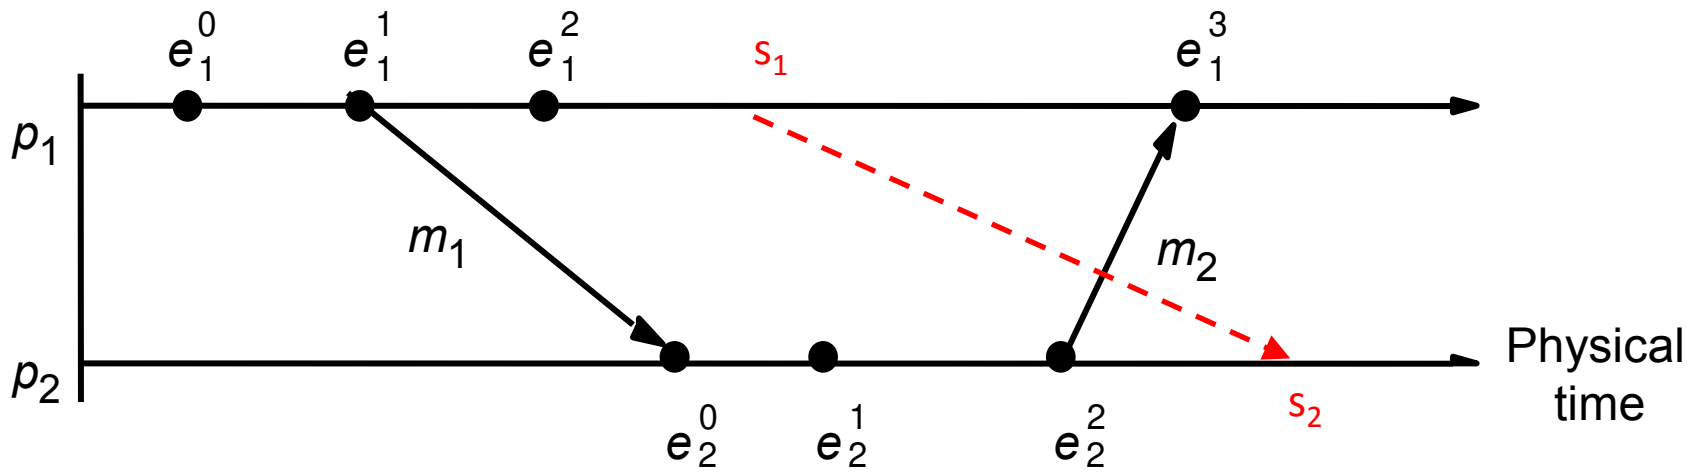
- *System model and assumptions:*
  - System of  $n$  processes:  $\langle p_1, p_2, p_3, \dots, p_n \rangle$ .
  - There are two uni-directional communication channels between each ordered process pair :  $p_j$  to  $p_i$  and  $p_i$  to  $p_j$ .
  - Communication channels are FIFO-ordered (first in first out).
  - All messages arrive intact, and are not duplicated.
  - No failures: neither channel nor processes fail.
- *Requirements:*
  - Snapshot should not interfere with normal application actions, and it should not require application to stop sending messages.
  - Any process may initiate algorithm.



# Chandy-Lamport Algorithm Intuition

- First, initiator  $p_i$ :
  - **records** its own state.
  - creates a special **marker** message.
  - sends the **marker** to all other process.
- When a process receives a **marker**.
  - **records** its own state.

# Chandy-Lamport Algorithm Intuition

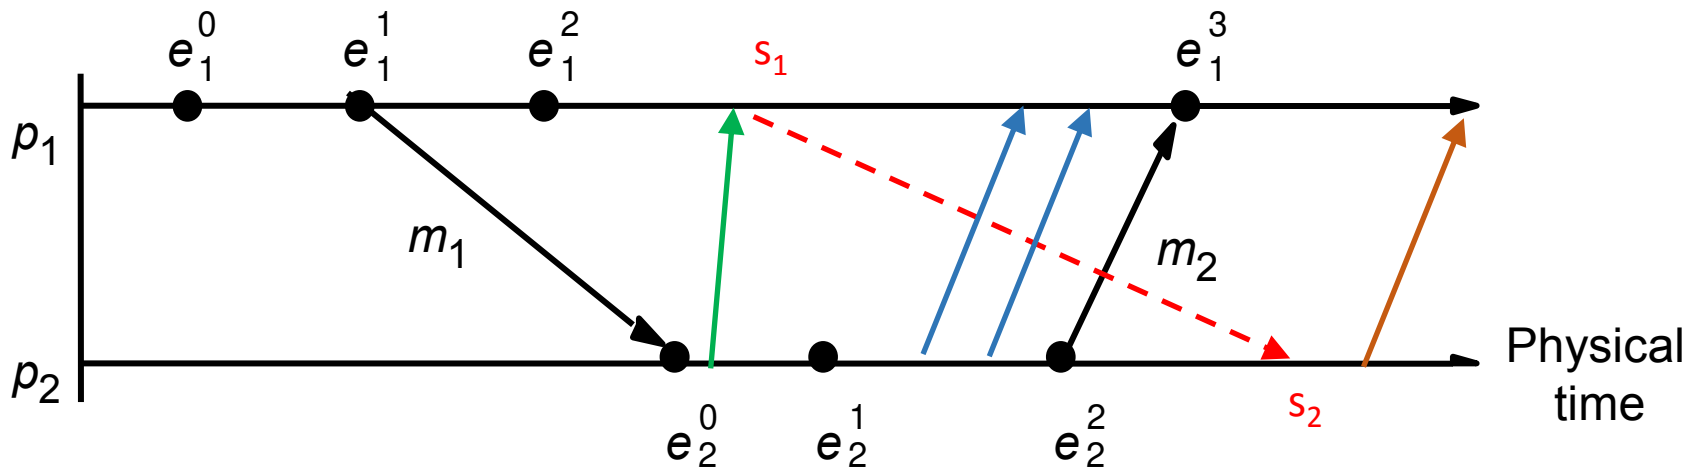


Cut frontier:  $\{e_1^2, e_2^2\}$

# Process, state, events

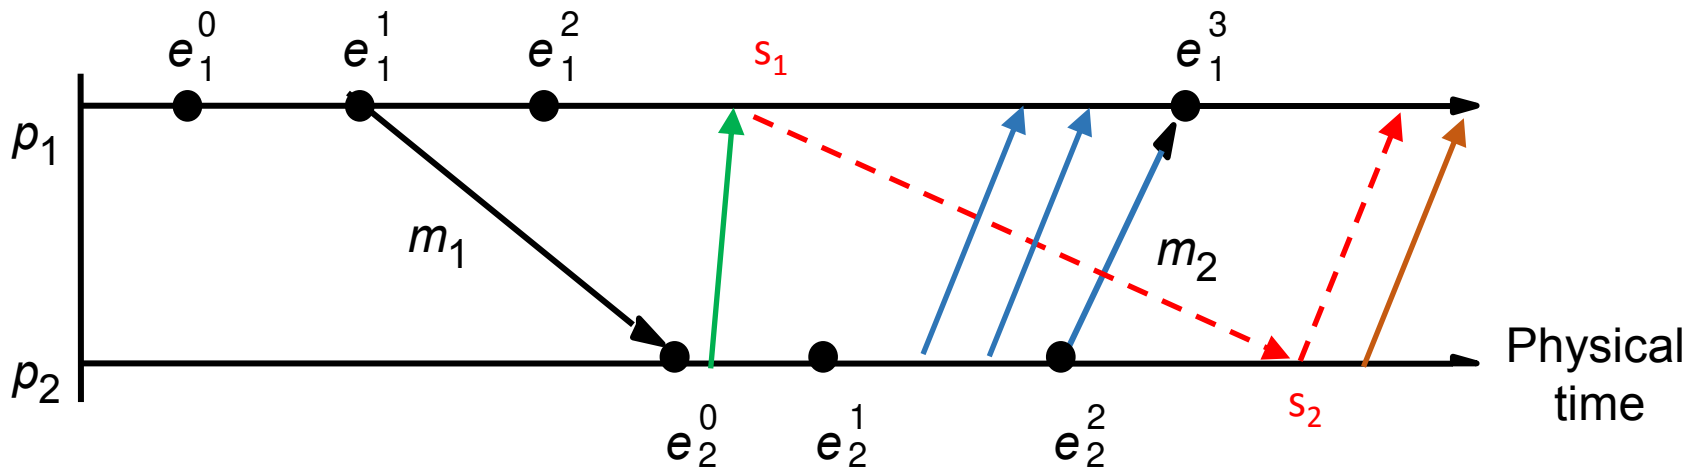
- Consider a system with  $n$  processes:  $\langle p_1, p_2, p_3, \dots, p_n \rangle$ .
- Each process  $p_i$  is associated with *state*  $s_i$ .
  - State includes values of all local variables, affected files, etc.
- Each channel can also be associated with a state.
  - Which messages are currently *pending* on the channel.
  - Can be computed from process' state:
    - Record when a process sends and receives messages.
    - if  $p_i$  sends a message that  $p_j$  has not yet received, it is pending on the channel.
- State of a process (or a channel) gets transformed when an *event* occurs. 3 types of events:
  - local computation, sending a message, receiving a message.

# Chandy-Lamport Algorithm Intuition



Cut frontier:  $\{e_1^2, e_2^2\}$

# Chandy-Lamport Algorithm Intuition



Cut frontier:  $\{e_1^2, e_2^2\}$

# Chandy-Lamport Algorithm Intuition

- First, initiator  $p_i$ :
  - **records** its own state.
  - creates a special **marker** message.
  - sends the **marker** to all other process.
  - start recording messages received on other channels.
    - until a marker is received on a channel.
- When a process receives a **marker**.
  - If marker is received for the first time.
    - **records** its own state.
    - sends **marker** on all other channels.
    - start recording messages received on other channels.
      - until a marker is received on a channel.